

Approximation-First Timeseries Query At Scale

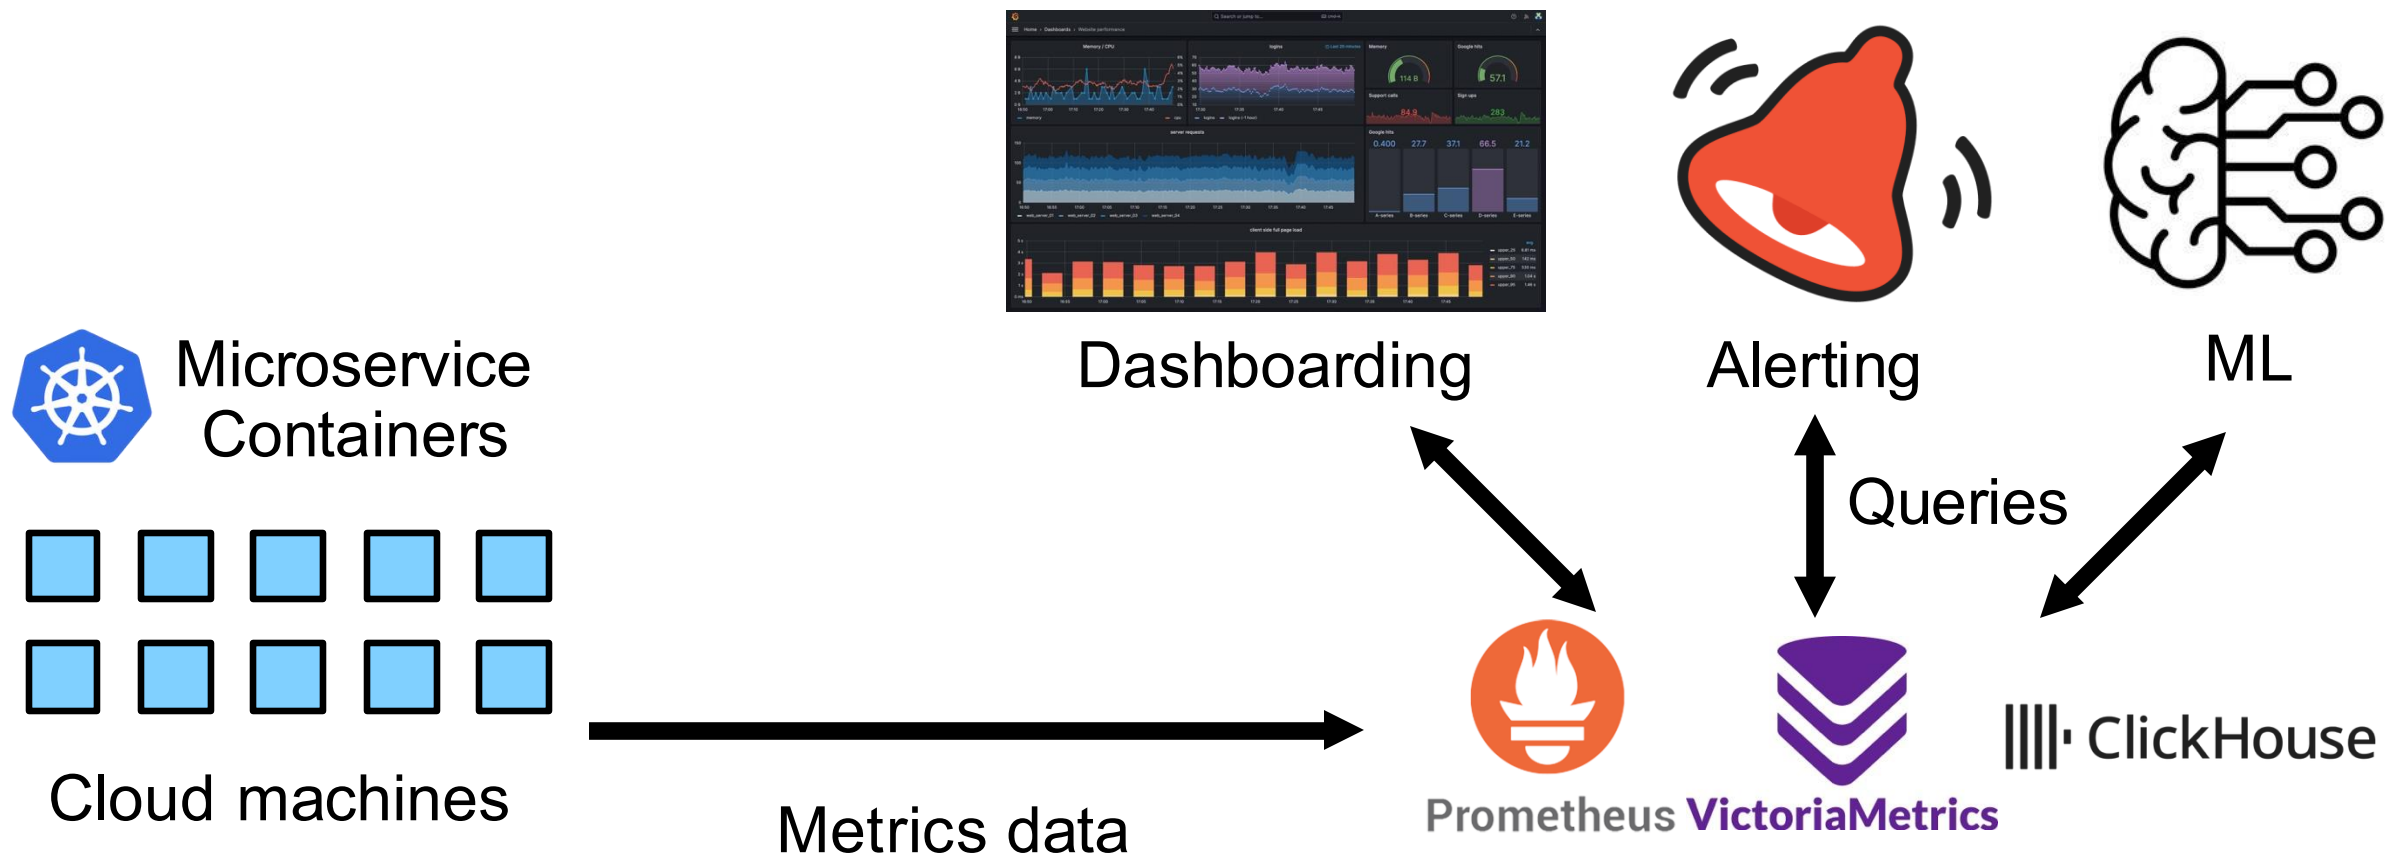
Zeying Zhu, Jonathan Chamberlain[†], Kenny Wu,

David Starobinski[†], Zaoxing Liu


University of Maryland, [†]Boston University



Need for Cloud and Network Observability



Example: Monitoring Cloud Machines



Cloud machines

CPU Usage /Timestamps	5:00:00	5:00:01	5:00:02	5:00:03	...
Host 1	10%	35%	15%	20%	
Host 2	20%	15%	19%	20%	
Host 3	15%	75%	80%	85%	
Host 4	90%	80%	95%	75%	
...					



Metrics data



Queries



Example *Rule* Queries for Cloud Resource Scaling

CPU Usage /Timestamps	5:00:00	...	5:01:00	...	5:02:00	...	5:03:00	...	5:04:00	...	5:05:00	...	5:06:00	...	5:07:00	...
Host 1	10%	...	35%	...	15%	...	20%	...	15%	...	20%	...	50%	...	15%	...
Host 2	20%	...	15%	...	19%	...	20%	...	19%	...	10%	...	40%	...	19%	...
Host 3	15%	...	75%	...	80%	...	85%	...	81%	...	80%	...	83%	...	80%	...
Host 4	90%	...	80%	...	95%	...	75%	...	85%	...	95%	...	80%	...	95%	...
...	Sliding Window Queries with overlaps between windows											

- *Rule* Queries: Periodically (every 1 min) querying 0.95-quantile of CPU usage of last 5 min, for each host.
 - `quantile_over_time(0.95, cpu_usage[5m])`

Bottlenecks in Window-based Queries

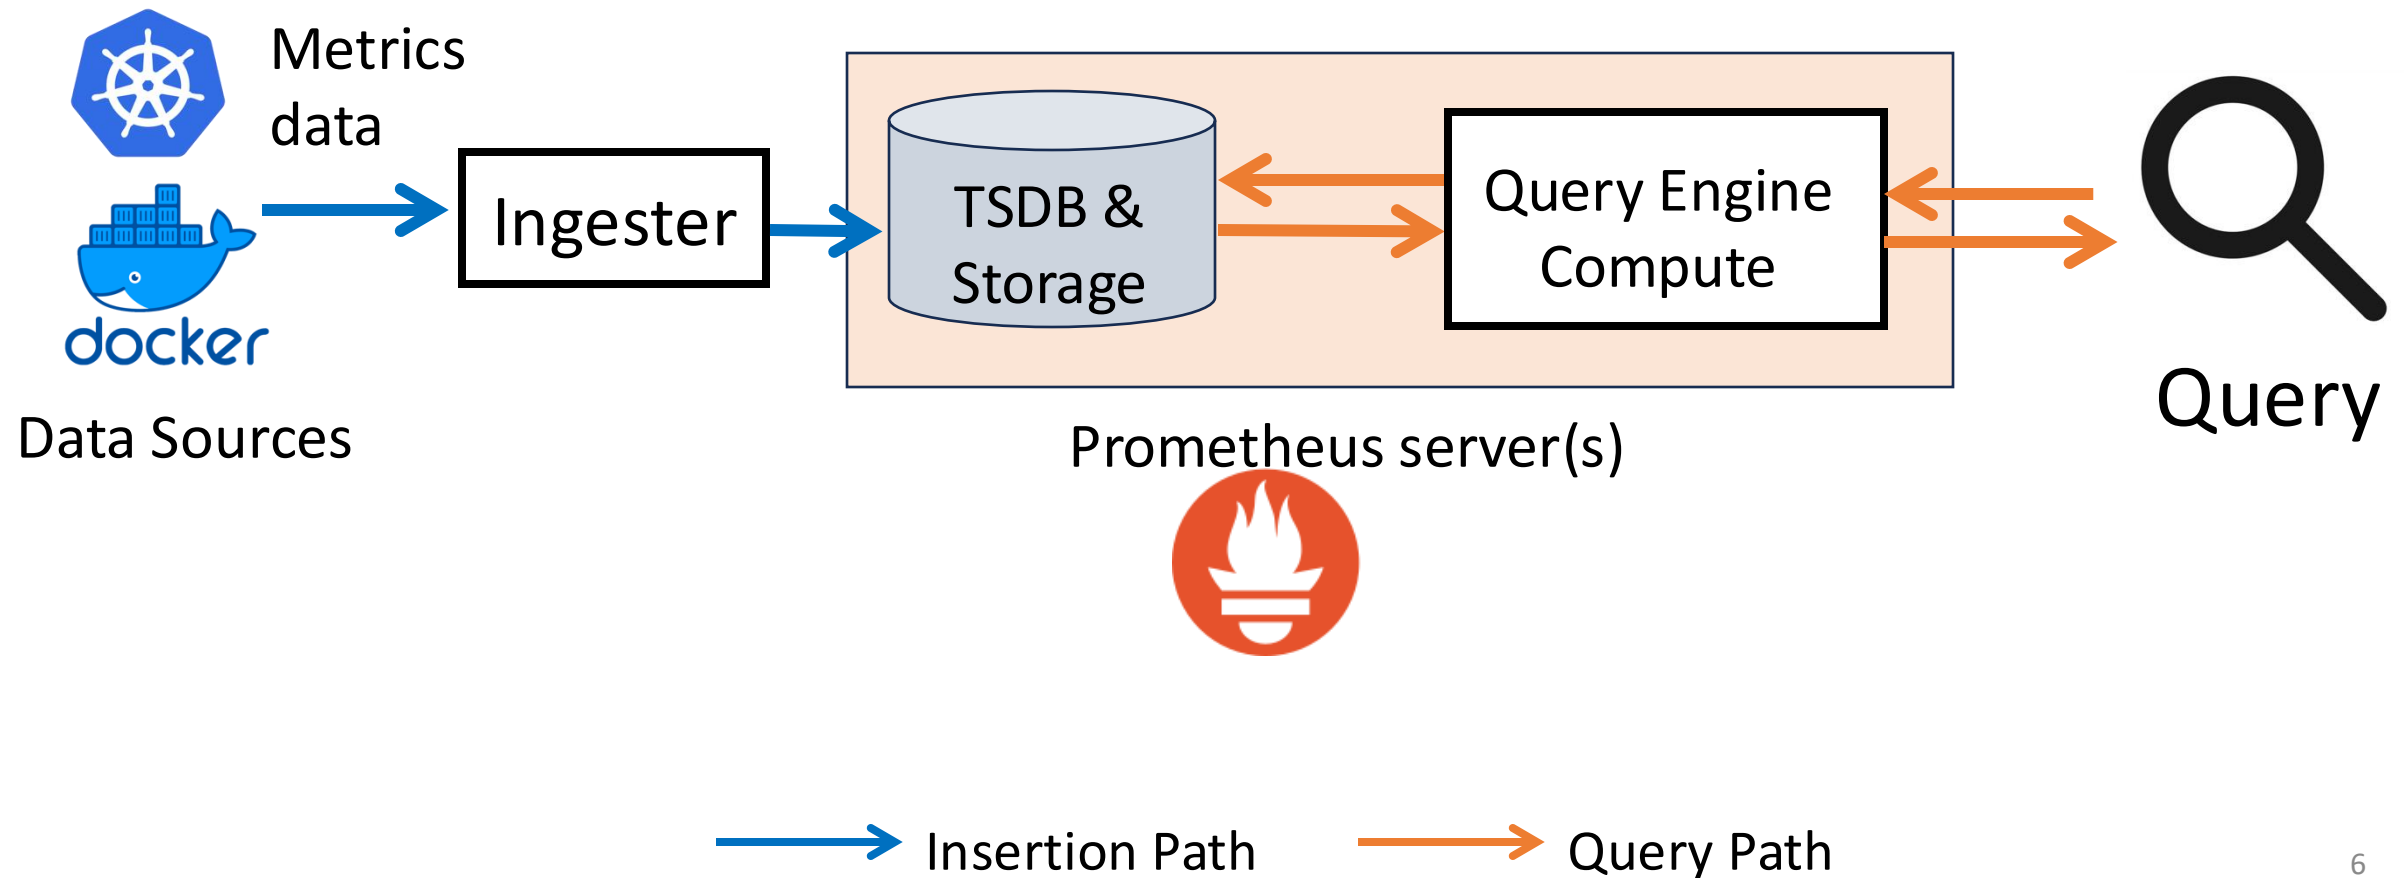
Table 2: CPU hotspots of evaluating a quantile rule query in Prometheus and VictoriaMetrics.

Func/Call Stack	CPU Time		Description
	Prometheus	VictoriaMetrics	
Data Scanning	41%	80.2%	Fetch data from storage
Query computation	27.6%	11.7%	Aggregation queries in rule
Go Garbage Collector	24.7%	4.3%	Golang garbage collector
mcall	4.5%	0.8%	Golang runtime scheduling

Major Bottlenecks due to window overlaps:

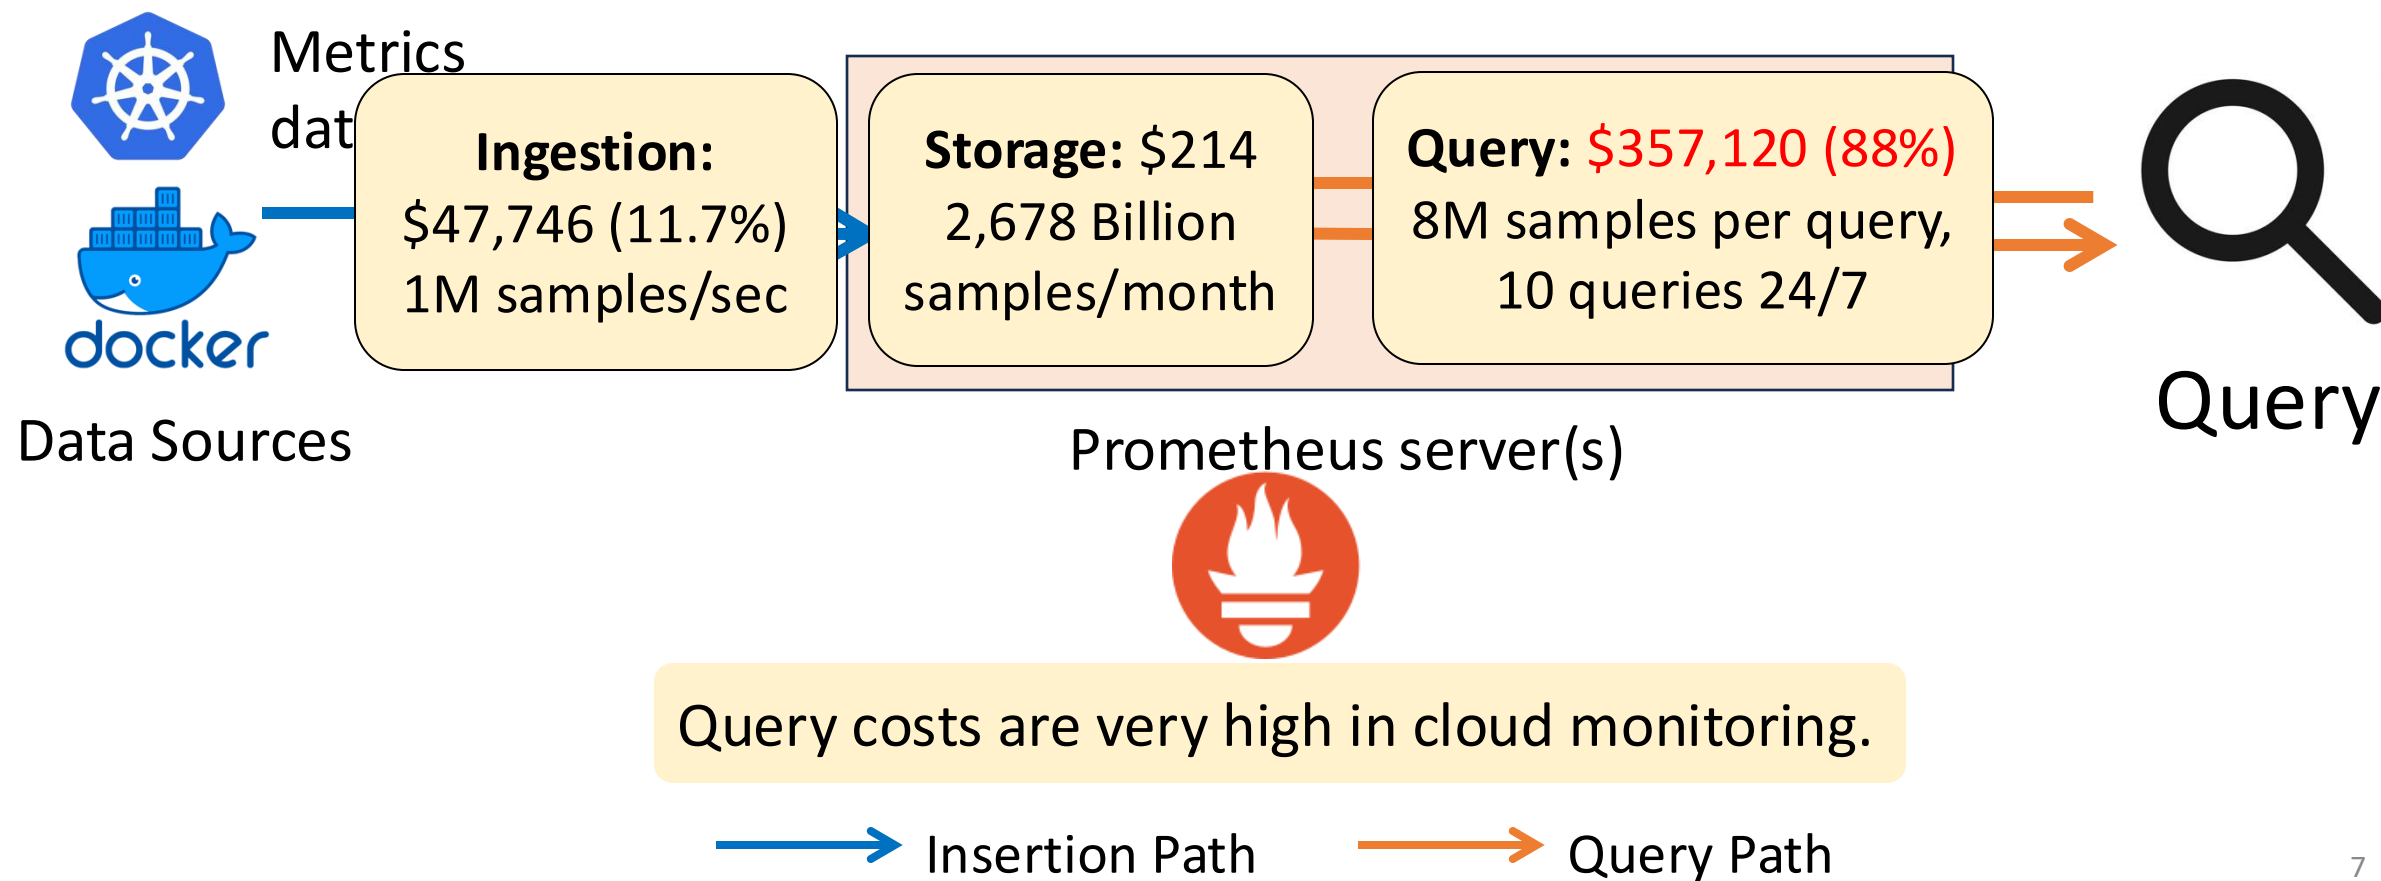
- Repeated data scanning from storage.
- Repeated and heavy query computation.

Prometheus is a Popular Observability Platform

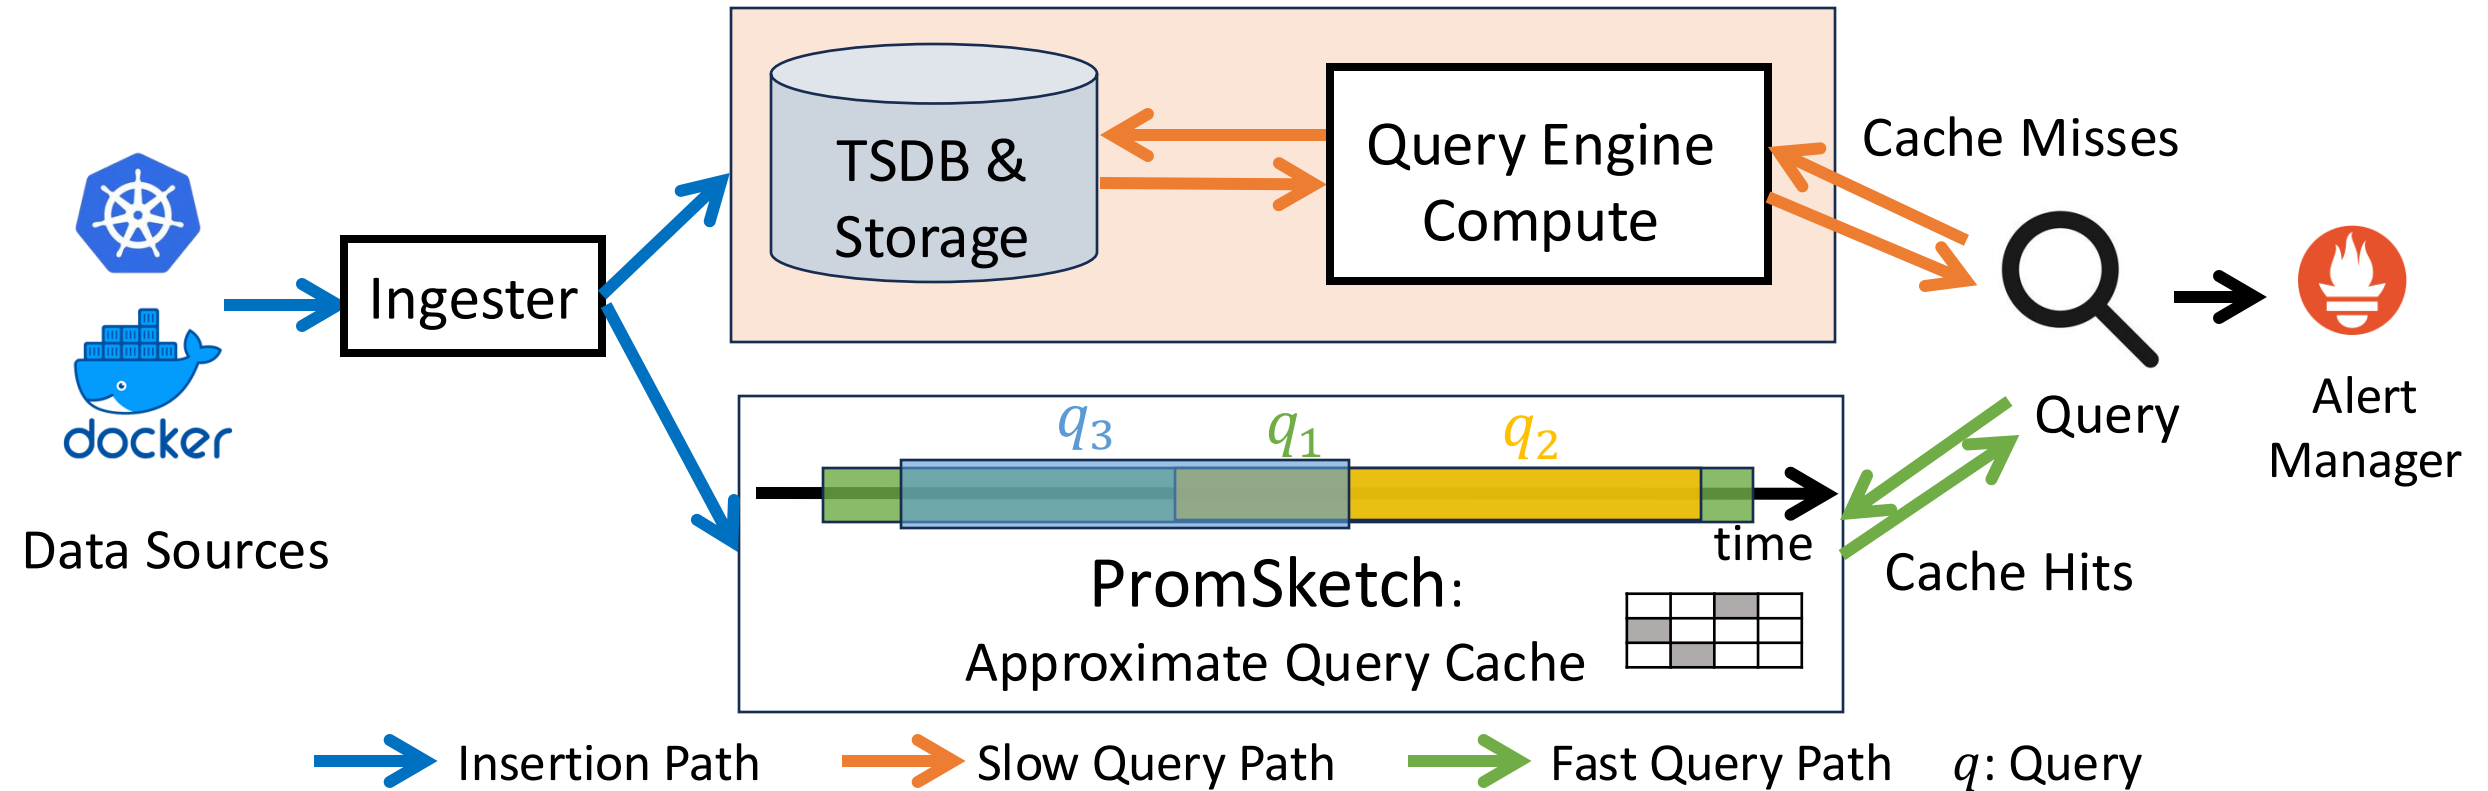


Cloud Monitoring Costs Example

Setup: monitoring 1000-node Kubernetes cluster, each node having 1000 metrics.
AWS Prometheus Pricing: Charge users based on the number of data samples processed.



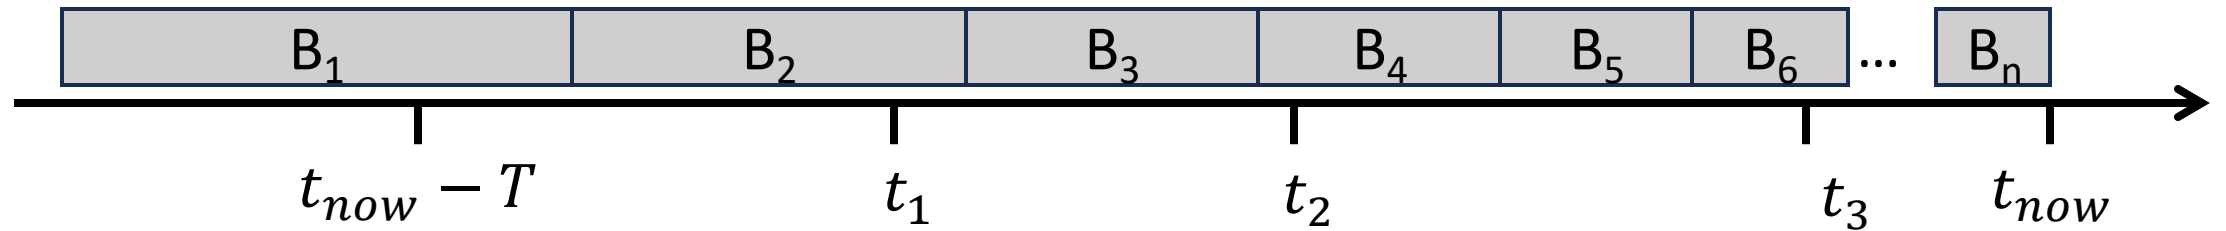
PromSketch as an Intermediate Cache



Key Idea: Combining arbitrary sub-window query frameworks and compact sketches as an intermediate query cache, computing overlapping windows once.

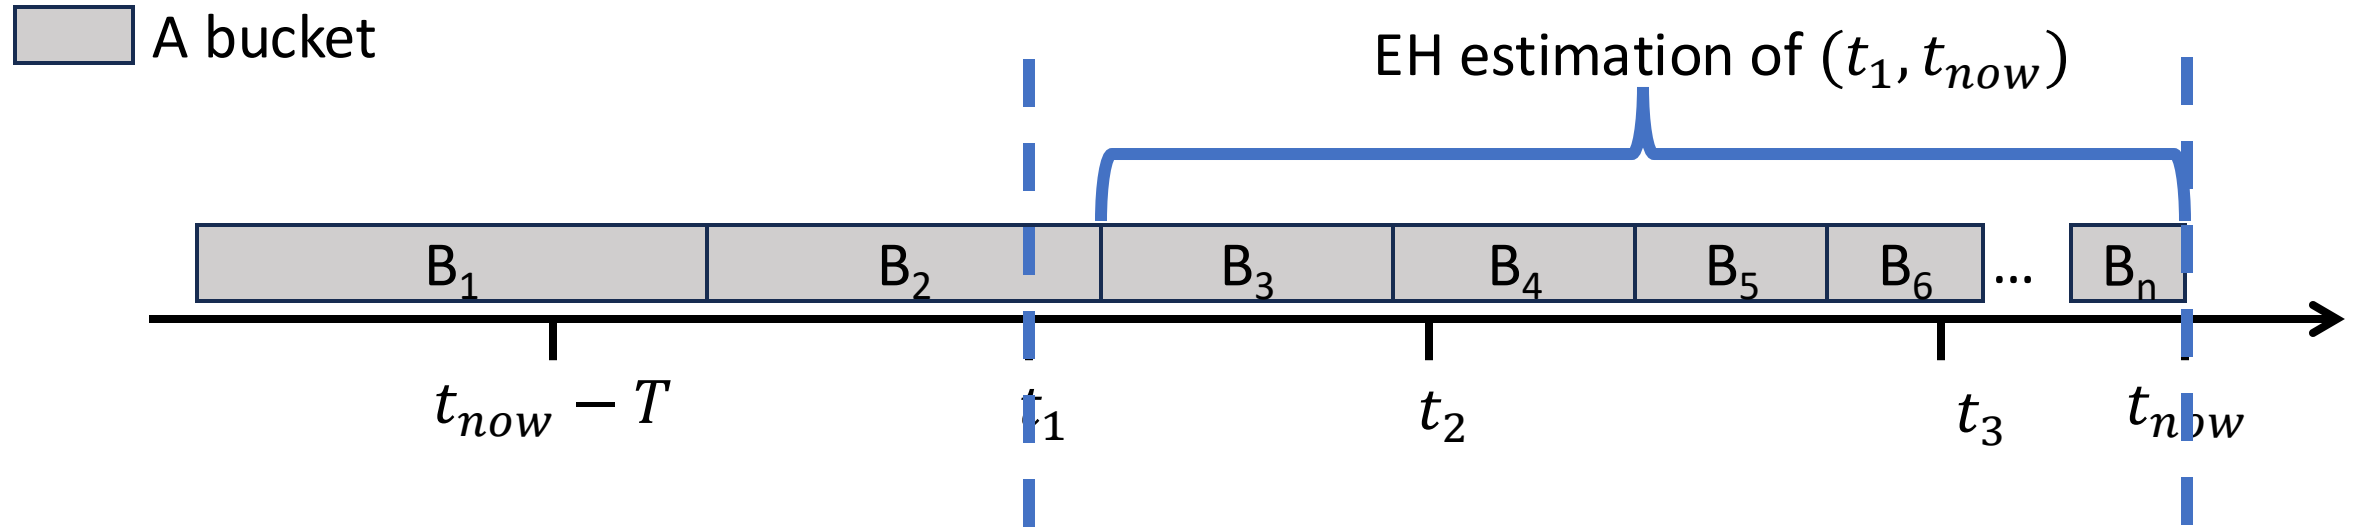
Exponential Histogram (EH) Framework

 A bucket



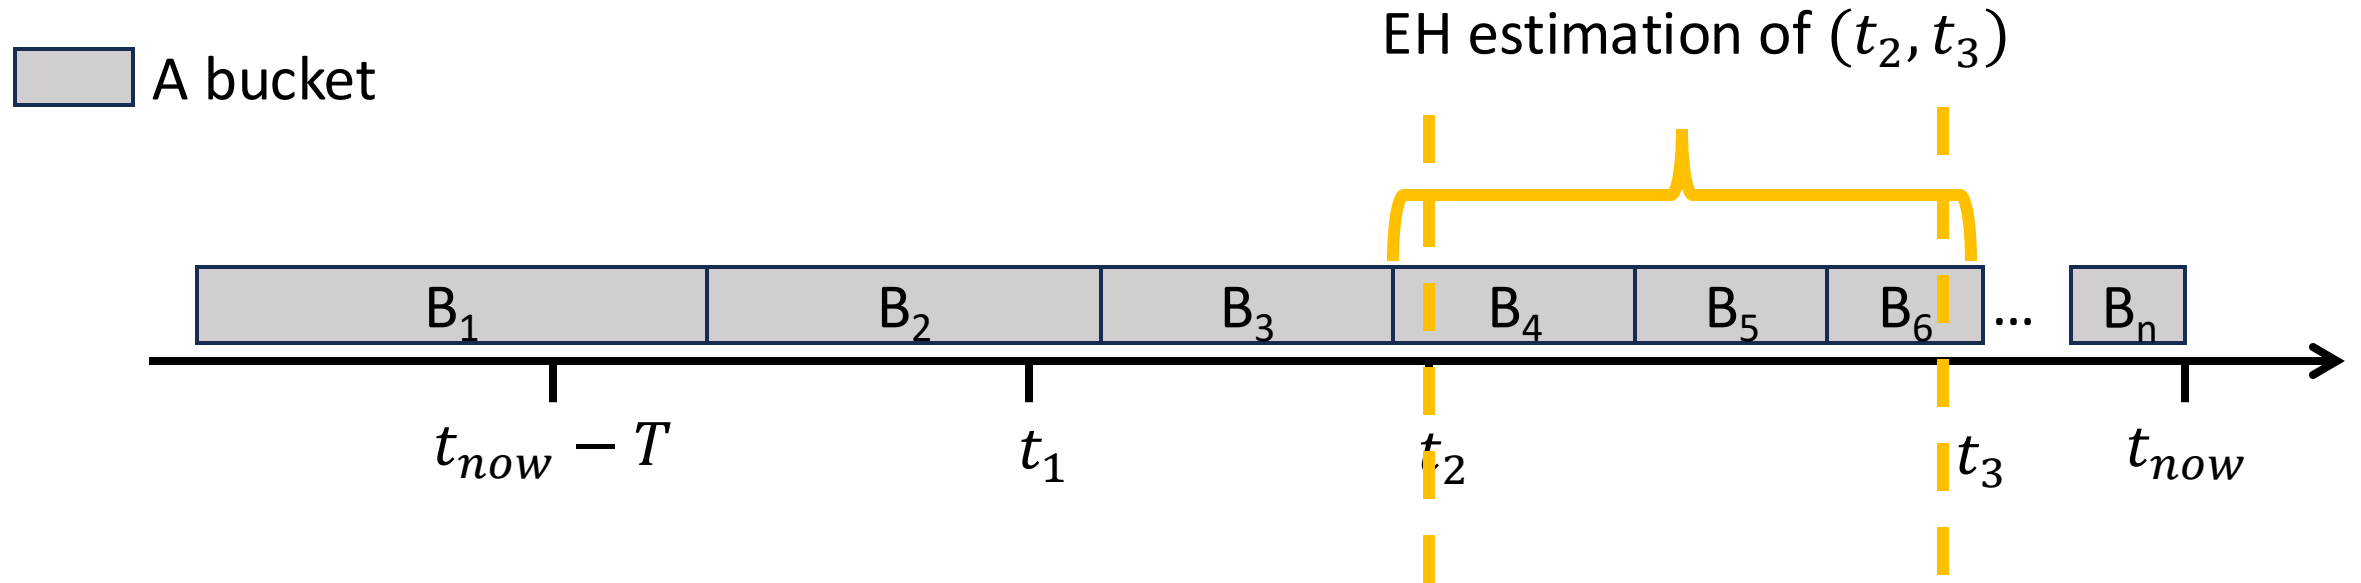
- EH supports querying arbitrary sub-windows cached by most recent T .

Querying Sub-windows by Merging Buckets



- EH supports querying arbitrary sub-windows cached by most recent T .
 - Combining buckets closest to the query time range.


Querying Sub-windows by Merging Buckets



- EH supports querying arbitrary sub-windows cached by most recent T .
 - Combining buckets closest to the query time range.

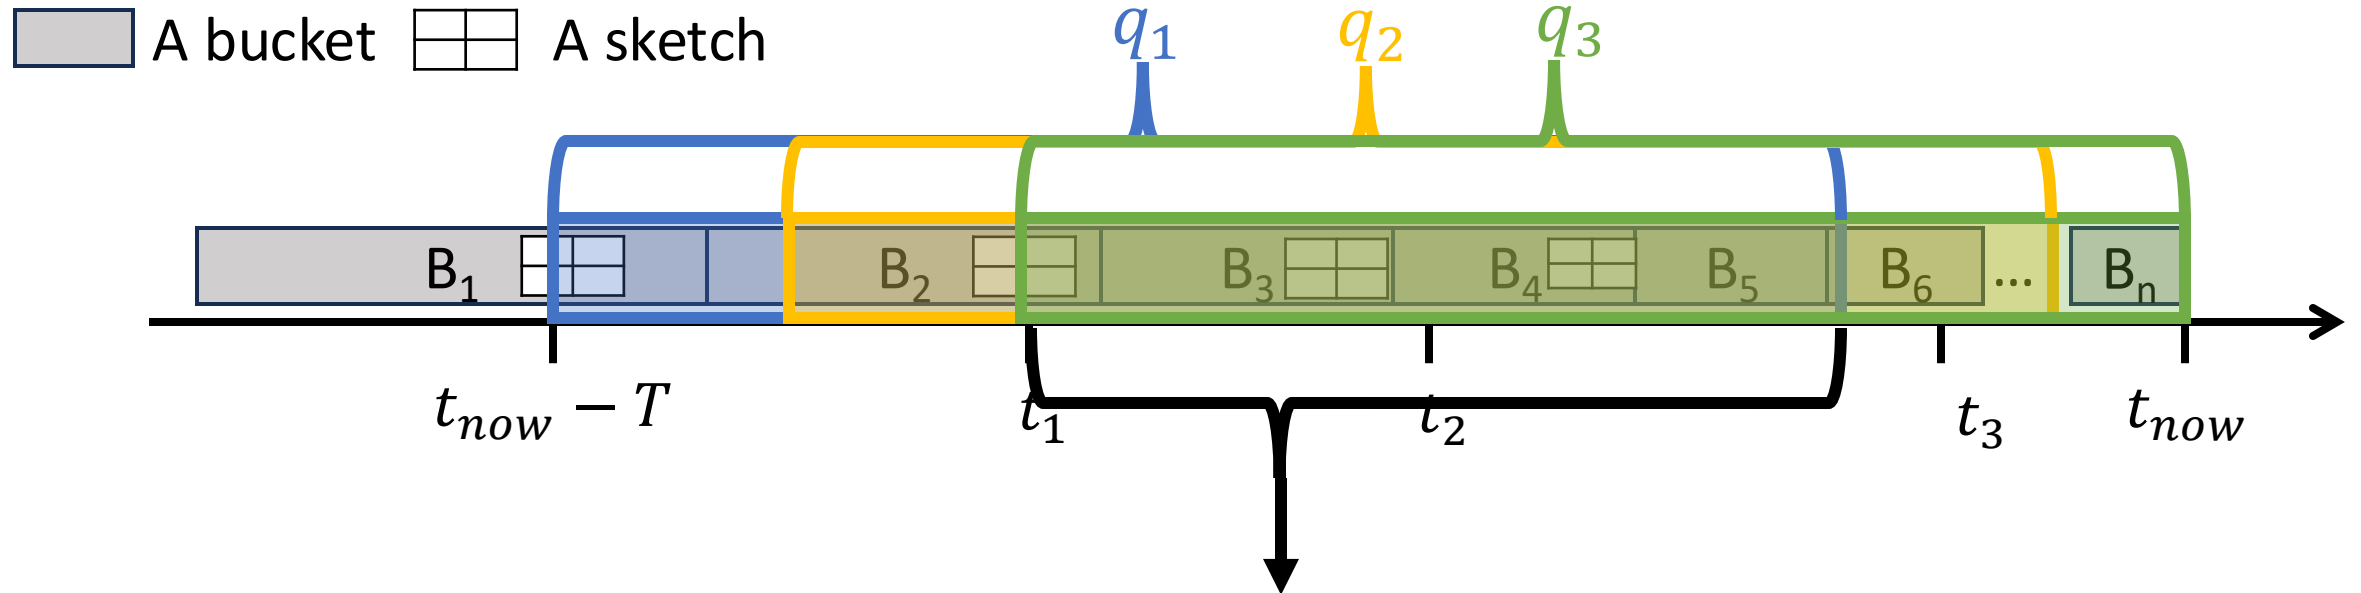
Provable Compact Sketches as Buckets

 A bucket  A sketch



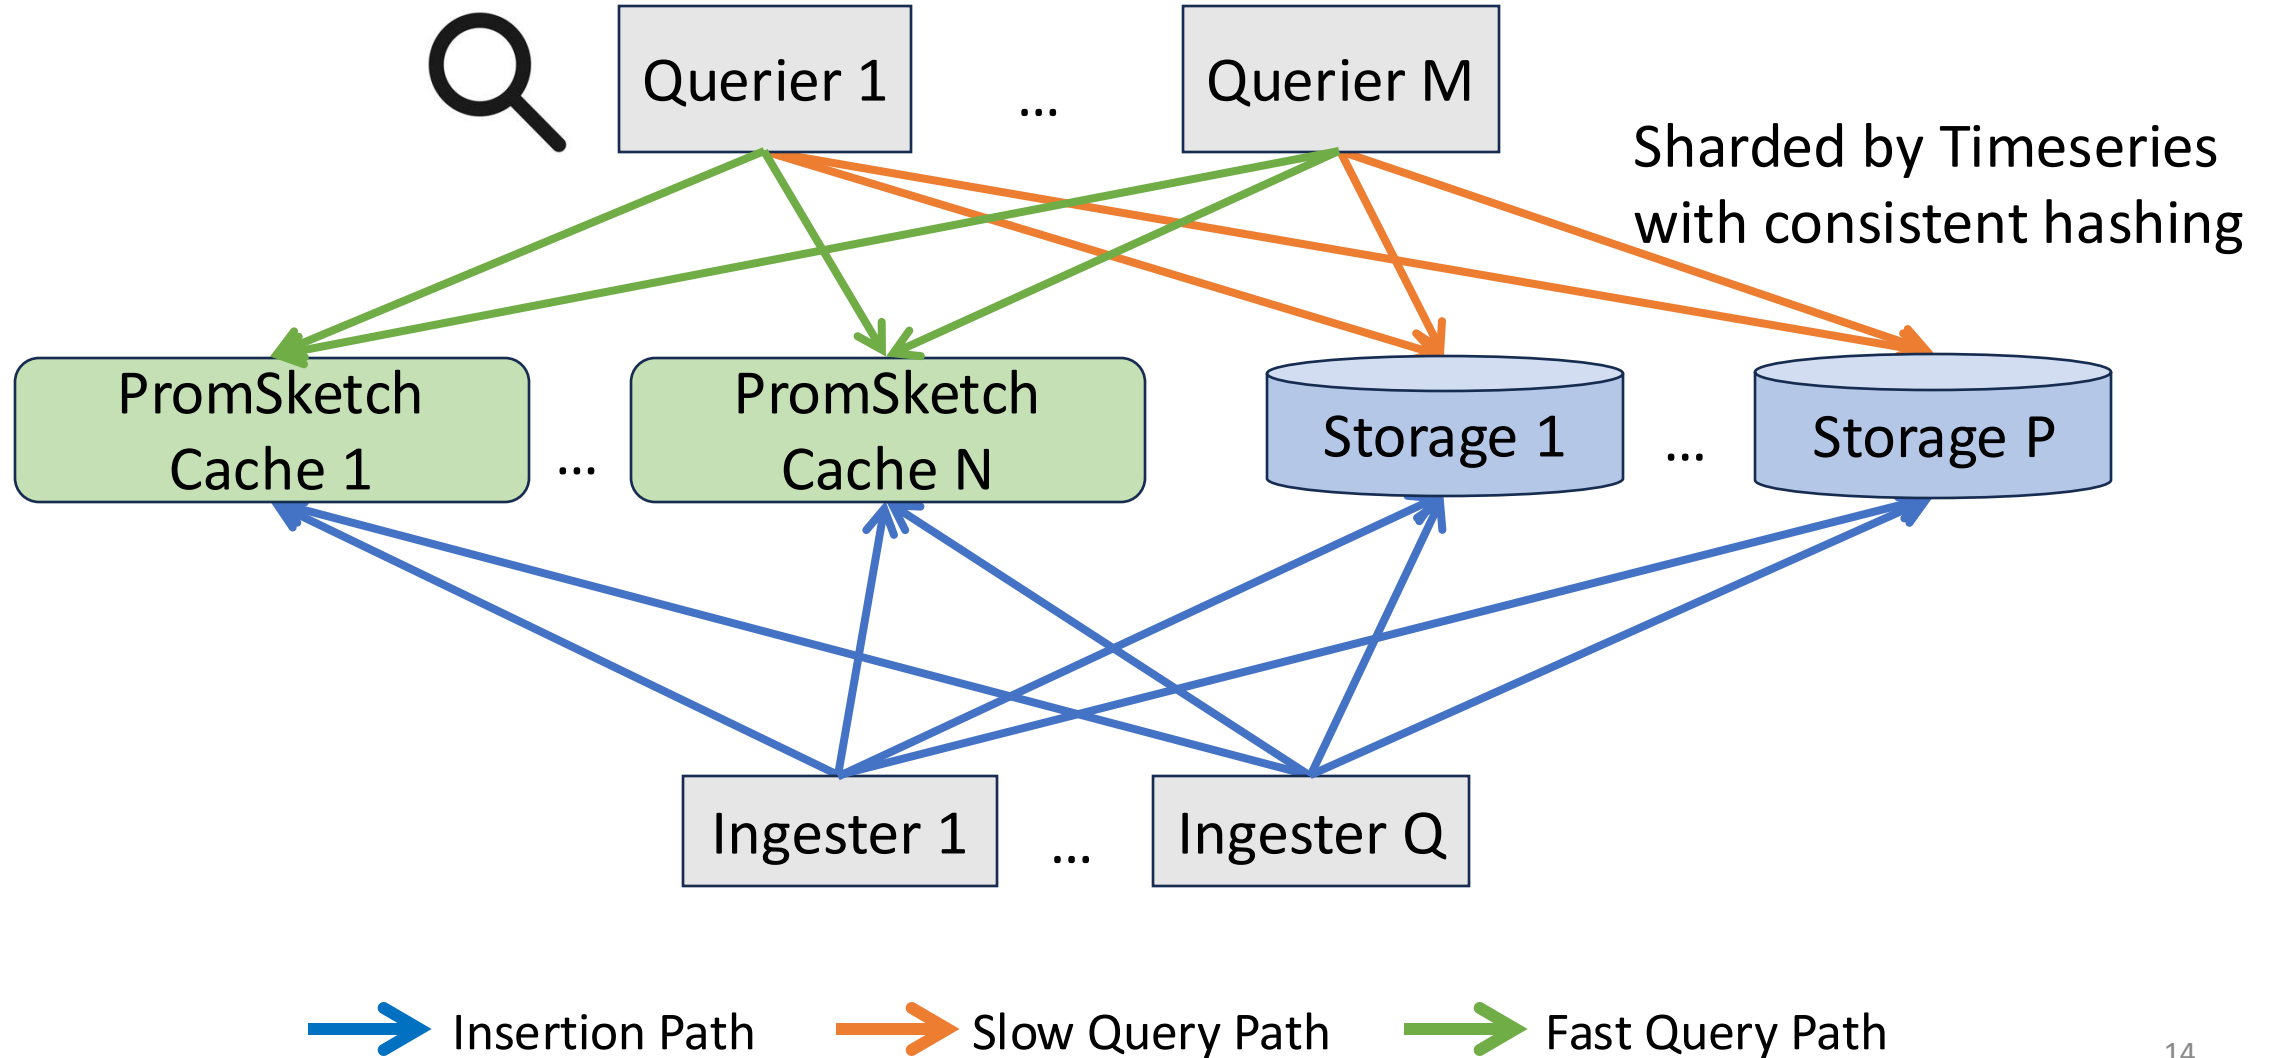
- KLL Sketch for Quantiles, Min, and Max.
- Universal Sketching for a wide range of functions, e.g., cardinality, L2 norm, entropy, top-K finding.

Precomputed Intermediate Cache Eliminates Query Overhead



- Overlapping windows are only computed once and cached in memory.

PromSketch Supports Cluster Version



Evaluation

- Testbed
 - Ubuntu servers with a 32-core CPU, 384GB DDR4 memory, and 1TB Seagate HD
- Baselines
 - Prometheus (single-machine)
 - Single-machine and cluster version VictoriaMetrics
 - Single-machine and distributed PromSketch
- Workloads
 - Google Cluster Dataset 2019 (Google), CAIDA NYU 2018 and CAIDA NYU 2019
 - Dynamic dataset with Zipf, uniform, normal distribution data samples
 - 10 million records


End-to-End Total Concurrent Query Latency

Total concurrent rule query latency on 10K-, 100K-, and 1M-sample windows

Metrics	Prometheus	Prometheus-PromSketch	VictoriaMetrics (VM)	VM-PromSketch
0.9-Quantile	5005 sec	28 sec (181x ↓)	96.1 sec	3.2 sec (30x ↓)
0.9-Quantile &Max &Average &Distinct	13177 sec	88.6 sec (154x ↓)	590 sec	9 sec (65x ↓)

- Up to 231x concurrent query latency reduction compared to Prometheus.
- Up to 158x compared to single-machine VictoriaMetrics.

Accuracy-Operational Cost Tradeoffs



- Achieve better cost-accuracy tradeoffs compared to uniform sampling.
- Achieve 5% mean errors with 5x~75x smaller cost compared to exact computation.

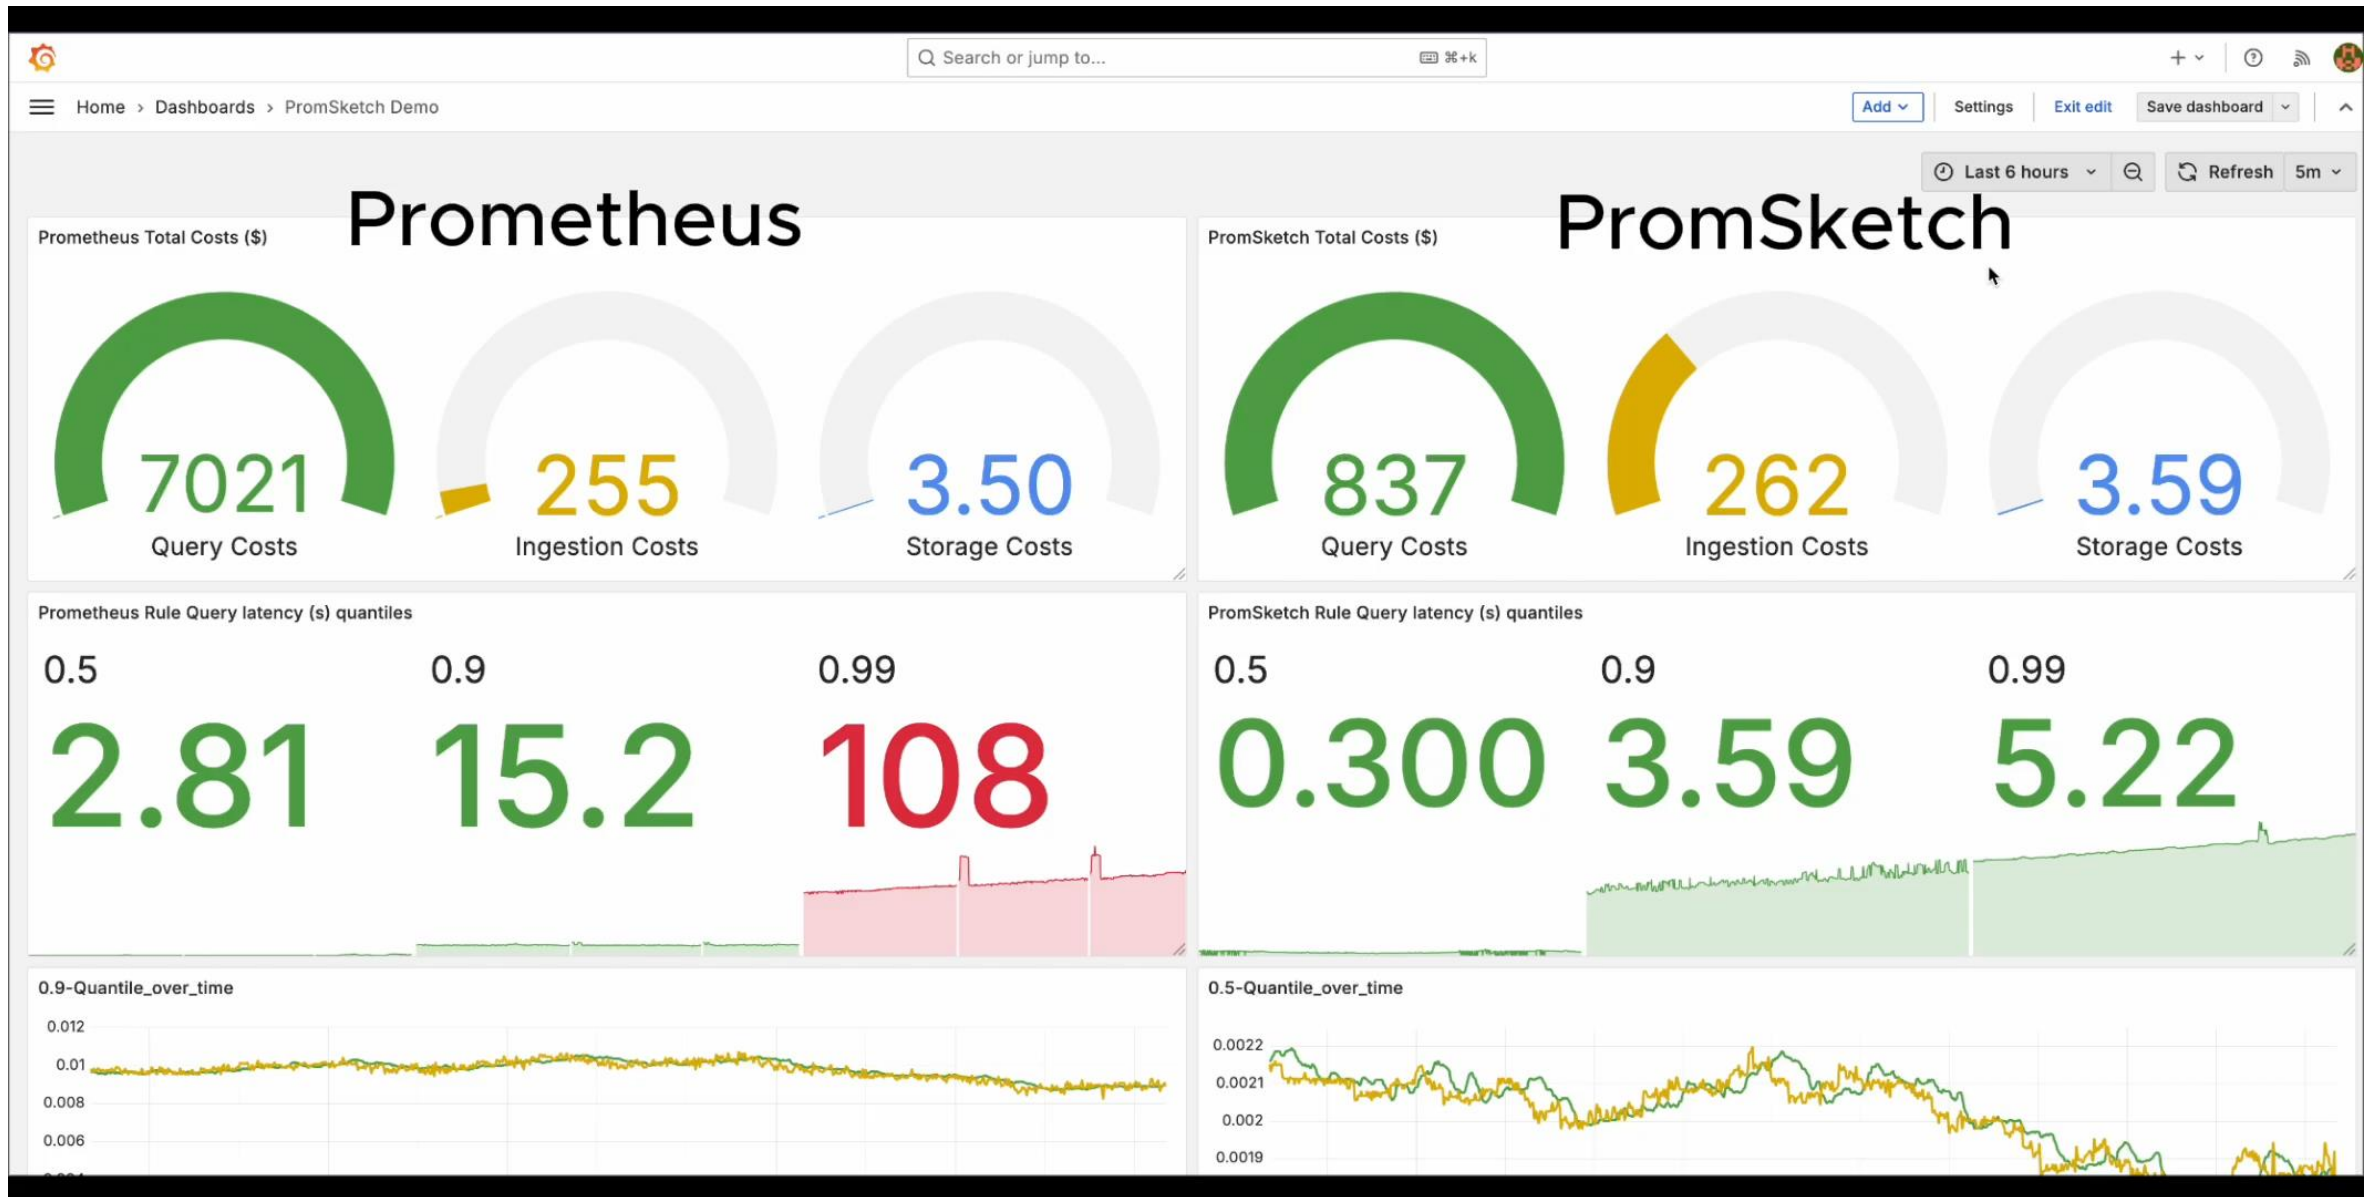
PromSketch Saves Monitoring Costs

Setup: monitoring 1000-node Kubernetes cluster, each node having 1000 metrics.

Resources	Data Volume	AWS Prometheus Costs/Month	Prometheus-PromSketch
Storage	2,678 Billion samples/month	\$214	\$214
Data Ingestion	1M samples/sec	\$47,746	\$47,746
Query Processing	8M samples per query, 10 queries running 24/7	\$357,120	\$267.8 (1,334x ↓)
Total Costs		\$405,080	\$48,227.8

- Reducing query costs by up to three orders of magnitude.

It's Demo Time!



PromSketch Conclusions

- Cloud timeseries monitoring systems, such as Prometheus, incur significant operational costs and high query latency.
- We design PromSketch, an approximation-first query processing framework that leverages sub-window query frameworks and sketch-based precomputation as intermediate caches.
- Integrated with Prometheus & VictoriaMetrics, open-sourced at <https://github.com/Froot-NetSys/promsketch>.

Thank you!